

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True)

```
1 import pandas as pd
2 df = pd.read_csv("/content/drive/MyDrive/Capstone/BigBasket Products.csv")
3 df.head()
```

	index	product	category	sub_category	brand	sale_price	market_pri
0	1	Garlic Oil - Vegetarian Capsule 500 mg	Beauty & Hygiene	Hair Care	Sri Sri Ayurveda	220.0	22
1	2	Water Bottle - Orange	Kitchen, Garden & Pets	Storage & Accessories	Mastercook	180.0	18

Next steps: [Generate code with df](#) [View recommended plots](#)

```
1 # Check data types of each column
2 df.info()
3
4 # Display basic statistics for numeric columns
5 df.describe()
6
7 # Display basic statistics for categorical columns
8 df.describe(include=['object'])
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 27555 entries, 0 to 27554

Data columns (total 10 columns):

#	Column	Non-Null Count	Dtype
0	index	27555 non-null	int64
1	product	27554 non-null	object
2	category	27555 non-null	object
3	sub_category	27555 non-null	object
4	brand	27554 non-null	object
5	sale_price	27555 non-null	float64
6	market_price	27555 non-null	float64
7	type	27555 non-null	object
8	rating	18929 non-null	float64
9	description	27440 non-null	object

dtypes: float64(3), int64(1), object(6)

memory usage: 2.1+ MB

	product	category	sub_category	brand	type	description
count	27554	27555	27555	27554	27555	27440
unique	23540	11	90	2313	426	21944
top	Turmeric Powder/Arisina Pudi	Beauty & Hygiene	Skin Care	Fresho	Face Care	A brand inspired by the Greek goddess of victo...

```
1 # Check for missing values
2 missing_values = df.isnull().sum()
3
4 # Display columns with missing values
5 missing_values[missing_values > 0]
```

```
product      1
brand         1
rating      8626
description  115
dtype: int64
```

```

1 # Numeric columns statistics
2 numeric_stats = df.describe()
3
4 # Categorical columns statistics
5 categorical_stats = df.describe(include=['object'])
6
7 numeric_stats, categorical_stats

```

↗

	count	index	sale_price	market_price	rating
count	27555.00000	27555.000000	27555.000000	27555.000000	18929.000000
mean	13778.00000	322.514808	382.056664	3.943410	
std	7954.58767	486.263116	581.730717	0.739063	
min	1.00000	2.450000	3.000000	1.000000	
25%	6889.50000	95.000000	100.000000	3.700000	
50%	13778.00000	190.000000	220.000000	4.100000	
75%	20666.50000	359.000000	425.000000	4.300000	
max	27555.00000	12500.000000	12500.000000	5.000000	

	product	category	sub_category	brand \
count	27554	27555	27555	27554
unique	23540	11	90	2313
top	Turmeric Powder/Arisina Pudi	Beauty & Hygiene	Skin Care	Fresho
freq	26	7867	2294	638

	type	description
count	27555	27440
unique	426	21944
top	Face Care	A brand inspired by the Greek goddess of victo...
freq	1508	47

```

1 # Checking for missing values again to confirm the columns and counts
2 missing_values = df.isnull().sum()
3 missing_values[missing_values > 0]
4
5 # Filling missing values for 'product' and 'brand' with mode
6 df['product'].fillna(df['product'].mode()[0], inplace=True)
7 df['brand'].fillna(df['brand'].mode()[0], inplace=True)
8
9 # Filling missing values for 'rating' with the mean
10 df['rating'].fillna(df['rating'].mean(), inplace=True)
11
12 # Filling missing values for 'description' with an empty string
13 df['description'].fillna('', inplace=True)
14
15 # Verify that there are no more missing values
16 df.isnull().sum()
17
18

```

↗

index	0
product	0
category	0
sub_category	0
brand	0
sale_price	0
market_price	0
type	0
rating	0
description	0

dtype: int64

```

1 import pandas as pd
2 import numpy as np
3
4 # Round up the 'rating' column to the nearest integer and convert to int
5 df['rating'] = df['rating'].apply(np.ceil).astype(int)
6
7 # Verify the change
8 print(df['rating'].head())
9
10

```

↗

0	5
1	3
2	4
3	4
4	5

Name: rating, dtype: int64

```

1 #Feature Selection
2 #For identifying key features influencing product ratings, we need to select relevant features from the dataset. Commonly con
3
4 # Select relevant features
5 selected_features = ['category', 'sub_category', 'brand', 'sale_price', 'market_price', 'type', 'description', 'rating']
6
7 # Create a new dataframe with the selected features
8 df_selected = df[selected_features]
9
10 # Display the first few rows of the new dataframe
11 df_selected.head()
12

```

	category	sub_category	brand	sale_price	market_price	type	descriptio
0	Beauty & Hygiene	Hair Care	Sri Sri Ayurveda	220.0	220.0	Hair Oil & Serum	This Produ contains Garl Oil that known
1	Kitchen, Garden & Pets	Storage & Accessories	Mastercook	180.0	180.0	Water & Fridge Bottles	Each produ is microwav safe (witho

Next steps:

[Generate code with df_selected](#)[View recommended plots](#)

```

1 #Data Preprocessing
2 #We'll now one-hot encode categorical variables (brand, category, sub_category) and prepare the dataset for further analysis.
3
4 # Selecting relevant features for analysis
5 selected_features = ['sale_price', 'brand', 'category', 'sub_category', 'rating']
6
7 # Create a new dataframe with selected features
8 df_selected = df[selected_features]
9
10 # One-hot encode categorical variables
11 df_encoded = pd.get_dummies(df_selected, columns=['brand', 'category', 'sub_category'], drop_first=True, dtype=int)
12
13 # Display the first few rows to verify
14 print(df_encoded.head())
15
16

```

→

	sale_price	rating	brand_&Stirred	brand_109°F	brand_137 Degree	\
0	220.0	5	0	0	0	
1	180.0	3	0	0	0	
2	119.0	4	0	0	0	
3	149.0	4	0	0	0	
4	162.0	5	0	0	0	
	brand_18 Herbs	brand_1mg	brand_1st Bites	brand_24 Mantra	brand_3 Roses	\
0	0	0	0	0	0	
1	0	0	0	0	0	
2	0	0	0	0	0	
3	0	0	0	0	0	
4	0	0	0	0	0	
	... sub_category_Skin Care	sub_category_Snacks & Namkeen	\			
0	...	0	0			
1	...	0	0			
2	...	0	0			
3	...	0	0			
4	...	0	0			
	sub_category_Snacks, Dry Fruits, Nuts	\				
0		0				
1		0				
2		0				
3		0				
4		0				
	sub_category_Spreads, Sauces, Ketchup	sub_category_Stationery	\			
0		0	0			
1		0	0			
2		0	0			
3		0	0			
4		0	0			
	sub_category_Steel Utensils	sub_category_Storage & Accessories	\			

0	0	0
1	0	1
2	0	0
3	0	0
4	0	0

	sub_category_Tea	sub_category_Tinned & Processed Food	sub_category_Water
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 2413 columns]

```

1 #Importing Libraries and Splitting Data and then implementing three classification algorithms: Decision Tree Classifier, Rand
2 #We'll evaluate each model's performance using cross-validation.
3
4 import pandas as pd
5 from sklearn.model_selection import train_test_split, cross_val_score
6 from sklearn.tree import DecisionTreeClassifier
7 from sklearn.ensemble import RandomForestClassifier
8 from sklearn.linear_model import LogisticRegression
9 from sklearn.metrics import classification_report, accuracy_score
10
11 # Selecting features and target
12 X = df_encoded.drop(columns=['rating'])
13 y = df_encoded['rating']
14
15 # Splitting the data into training and testing sets (80% train, 20% test)
16 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
17
18 # Display the shapes of the train and test sets to verify
19 print("Training set shape:", X_train.shape, y_train.shape)
20 print("Testing set shape:", X_test.shape, y_test.shape)
21

```

↗ Training set shape: (22044, 2412) (22044,)
 Testing set shape: (5511, 2412) (5511,)

```

1 #Implementing Classification Algorithms
2
3 from sklearn.model_selection import cross_val_score
4 from sklearn.tree import DecisionTreeClassifier
5 from sklearn.ensemble import RandomForestClassifier
6 from sklearn.metrics import classification_report, accuracy_score
7
8 # Initialize Decision Tree and Random Forest Classifiers
9 dt_classifier = DecisionTreeClassifier(random_state=42)
10 rf_classifier = RandomForestClassifier(random_state=42)
11
12 # List of classifiers
13 classifiers = [('Decision Tree', dt_classifier),
14               ('Random Forest', rf_classifier)]
15
16 # Evaluate each classifier using cross-validation and on the test set
17 for clf_name, clf in classifiers:
18     print(f"Training and evaluating {clf_name}...")
19
20     # Cross-validation
21     scores = cross_val_score(clf, X_train, y_train, cv=5, scoring='accuracy', n_jobs=-1)
22     print(f"Cross-validation scores: {scores}")
23     print(f"Mean accuracy: {scores.mean():.3f}")
24     print()
25
26     # Fit the classifier on the training data
27     clf.fit(X_train, y_train)
28
29     # Predict on the test data and evaluate
30     y_pred = clf.predict(X_test)
31     accuracy = accuracy_score(y_test, y_pred)
32     print(f"{clf_name} Accuracy on test set: {accuracy:.3f}")
33     print(classification_report(y_test, y_pred))
34     print()

```

↗ Training and evaluating Decision Tree...
 Cross-validation scores: [0.61193014 0.62463144 0.61510547 0.61555908 0.61297641]

Mean accuracy: 0.616

Decision Tree		Accuracy on test set: 0.625			
	precision	recall	f1-score	support	
1	0.15	0.16	0.15	73	
2	0.06	0.06	0.06	72	
3	0.16	0.17	0.16	252	
4	0.69	0.73	0.71	3065	
5	0.62	0.57	0.59	2049	
accuracy			0.62	5511	
macro avg			0.33	5511	
weighted avg			0.62	5511	

Training and evaluating Random Forest...

Cross-validation scores: [0.64118848 0.64254933 0.63461102 0.64073486 0.63112523]

Mean accuracy: 0.638

Random Forest		Accuracy on test set: 0.648			
	precision	recall	f1-score	support	
1	0.13	0.10	0.11	73	
2	0.08	0.06	0.06	72	
3	0.19	0.16	0.18	252	
4	0.70	0.75	0.72	3065	
5	0.64	0.60	0.62	2049	
accuracy			0.65	5511	
macro avg			0.35	5511	
weighted avg			0.64	5511	

```

1 from sklearn.model_selection import train_test_split, cross_val_score
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.metrics import classification_report, accuracy_score
4 from sklearn.preprocessing import OneHotEncoder
5 from sklearn.compose import ColumnTransformer
6 from sklearn.pipeline import Pipeline
7 import pandas as pd
8
9 # Splitting data into features (X) and target variable (y)
10 X = df.drop(columns=['rating']) # Features
11 y = df['rating'] # Target variable
12
13 # Identify categorical columns
14 categorical_columns = X.select_dtypes(include=['object']).columns
15
16 # Create a column transformer with OneHotEncoder for categorical columns
17 preprocessor = ColumnTransformer(
18     transformers=[
19         ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_columns)
20     ],
21     remainder='passthrough'
22 )
23
24 # Create a pipeline with the preprocessor and Logistic Regression classifier
25 logreg_pipeline = Pipeline(steps=[
26     ('preprocessor', preprocessor),
27     ('classifier', LogisticRegression(max_iter=1000, random_state=42))
28 ])
29
30 # Splitting data into training and testing sets
31 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
32
33 # Train and evaluate Logistic Regression
34 print("Training and evaluating Logistic Regression...")
35 logreg_pipeline.fit(X_train, y_train)
36
37 # Cross-validation
38 print("Cross-validation scores:")
39 cv_scores = cross_val_score(logreg_pipeline, X_train, y_train, cv=5, scoring='accuracy', n_jobs=-1)
40 print(cv_scores)
41 print(f"Mean accuracy: {cv_scores.mean():.3f}")
42 print()
43
44 # Evaluate on the test set
45 y_pred_logreg = logreg_pipeline.predict(X_test)
46 accuracy_logreg = accuracy_score(y_test, y_pred_logreg)
47 print(f"Logistic Regression Accuracy on test set: {accuracy_logreg:.3f}")
48 print(classification_report(y_test, y_pred_logreg))
49
50
51

```

→ Training and evaluating Logistic Regression...
 Cross-validation scores:
 /usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```

n_iter_i = _check_optimize_result(
[0.59423906 0.55069177 0.56112497 0.55091858 0.59255898]
Mean accuracy: 0.570

```

Logistic Regression Accuracy on test set: 0.556

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

1	0.00	0.00	0.00	73
2	0.20	0.01	0.03	72
3	1.00	0.00	0.01	252
4	0.56	1.00	0.71	3065
5	0.50	0.00	0.00	2049

accuracy			0.56	5511
macro avg	0.45	0.20	0.15	5511
weighted avg	0.54	0.56	0.40	5511

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1344: UndefinedMetricWarning: Precision and F-score

```

warn prf(average. modifier. msd start. len(result))
1 from sklearn.metrics import confusion_matrix
2
3 # Initialize Decision Tree and Random Forest Classifiers
4 dt_classifier = DecisionTreeClassifier(random_state=42)
5 rf_classifier = RandomForestClassifier(random_state=42)
6 logreg_classifier = LogisticRegression(max_iter=1000, random_state=42)
7
8 # List of classifiers
9 classifiers = [('Decision Tree', dt_classifier),
10              ('Random Forest', rf_classifier),
11              ('Logistic Regression', logreg_classifier)]
12
13 # Evaluate each classifier and print confusion matrix
14 for clf_name, clf in classifiers:
15     print(f"Training and evaluating {clf_name}...")
16
17     # Fit the classifier on the training data
18     clf.fit(X_train, y_train)
19
20     # Predict on the test data
21     y_pred = clf.predict(X_test)
22
23     # Calculate confusion matrix
24     cm = confusion_matrix(y_test, y_pred)
25
26     # Print confusion matrix
27     print(f"Confusion Matrix for {clf_name}:")
28     print(cm)
29     print()
30
31

```

↩ Training and evaluating Decision Tree...

Confusion Matrix for Decision Tree:

```

[[ 12   5   6  33  17]
 [   1   4   6  43  18]
 [   8   9  42 140  53]
 [  42  35 145 225 618]
 [  19  13  69 789 1159]]

```

Training and evaluating Random Forest...

Confusion Matrix for Random Forest:

```

[[   7   4   6  39  17]
 [   0   4   7  43  18]
 [   4   6  41 147  54]
 [  30  27 106 229 612]
 [  13  11  53 742 1230]]

```

Training and evaluating Logistic Regression...

/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>